

# Digital Afterlife AI (Griefbots)

*Ethical Challenges of AI Systems that Simulate Deceased Individuals*

**Name:** Chris Kevin A

**Register No.:** 2460349

# Background of the Case

02

## Organizations / Systems

- Project December (Jason Rohrer, 2021) – GPT-3-powered chatbot recreating deceased personas
- HereAfter AI (James Vlahos, 2017) – voice-based "Dadbot" preserving parent memories
- StoryFile (2020) – interactive video AI for post-death communication
- Microsoft patent (US 2021/0004455) – technology to create chatbots from deceased users' data

## Purpose of the AI System

- Provide emotional comfort and grief support to the bereaved
- Preserve memories, personality, and communication style of the deceased
- Enable ongoing 'conversations' using text, voice, and video synthesis
- Digitally immortalise individuals through their digital footprint

## Timeline of Key Events

- 2017: James Vlahos builds HereAfter AI after father's cancer diagnosis
- 2020: StoryFile launches; used at David Galen's memorial service
- 2021: Project December publicly enables griefbot creation via GPT-3
- 2021: Microsoft granted patent for personal chatbot from deceased data
- 2023: Meta AI and Character.AI offer user-customised persona simulation

## What Happened?

Digital Afterlife AI systems collect a deceased person's messages, emails, voice recordings, social media posts, and photos to train personalised AI models.

Platforms such as Project December enabled users to "converse" with AI simulations of dead loved ones, sometimes without the knowledge or consent of the deceased or their next of kin.

A 2021 MIT Technology Review case documented a woman using Project December to chat with an AI trained on texts from her deceased fiancé, raising widespread ethical debate.

## Key Ethical Issues

- Posthumous data use without prior or informed consent
- Psychological manipulation of grieving individuals
- Identity appropriation and misrepresentation of deceased
- Lack of regulatory governance and transparency

## Stakeholders Affected

Bereaved Families

Deceased Individuals

Platform Developers

Regulatory Bodies

Mental Health Community

General Public

**Key Stat:** A 2023 survey by YouGov found that 39% of Americans would consider using a digital afterlife service to communicate with a deceased loved one, yet fewer than 5% of those users had formally consented to such use of their own data (YouGov, 2023).

## AI Technology Used

- Large Language Models (GPT-3/4) for text generation and personality mimicry
- Natural Language Processing (NLP) for tone, vocabulary and speech pattern replication
- Voice cloning via WaveNet / Tacotron for audio simulation
- Deepfake video synthesis (e.g., StoryFile hologram-style videos)
- Emotion modelling through affective computing algorithms

## Data Sources

- Text messages, emails, and chat logs provided by family
- Social media content: posts, comments, captions (Facebook, Twitter, Instagram)
- Voice recordings: calls, videos, podcasts, interviews
- Photographs used for visual synthesis and deepfake generation
- Personal documents, diaries, and biographical materials

## Decision-Making Process

- Contextual retrieval: user query matched to closest training example
- Probabilistic generation: model produces responses consistent with persona patterns
- Emotional tone calibration based on grieving context
- No real-time human oversight during live interactions
- Opaque model weights – families cannot audit what the AI 'knows'

## Benefits (Who gains?)

- Bereaved individuals experience reduced immediate grief and loneliness
- Preservation of family heritage and cultural memory across generations
- Therapeutic potential: some grief counsellors report short-term emotional relief
- Historical value: preserving voices of significant public figures (e.g., StoryFile Holocaust survivor project)

## Harms (Who is harmed?)

- Prolonged or complicated grief: unnatural mourning delays acceptance (APA, 2022)
- Deceased individuals' autonomy violated – data used without ante-mortem consent
- Vulnerable griever manipulated by commercially incentivised platforms
- Normalisation of death commodification and psychological dependency

## Duties & Obligations Violated

- Kant's Categorical Imperative demands that individuals never be used merely as means. Using a deceased person's data to generate a commercial product treats them as a means without their rational consent.
- Duty of honesty: Griefbots simulate responses the deceased never made. This constitutes systematic deception, violating the deontological duty of truthfulness.
- Breach of the duty to protect personal data: Most jurisdictions do not extend GDPR-style data rights post-mortem, creating a legal and moral gap.
- Fiduciary duty violation: Platforms collecting user data implicitly commit to use it only as originally intended – not for posthumous AI training.

## Fairness & Transparency Assessment

- Lack of meaningful informed consent: Users cannot meaningfully consent to future AI training that did not exist when data was collected.
- Opaque algorithms: The deceased's 'personality' is a probabilistic model – there is no transparency about how responses are generated or where errors occur.
- Asymmetric power: Platforms hold all data control; families and estates have minimal rights or recourse under current law.
- No standardised disclosure: Users interacting with griefbots are rarely clearly informed they are talking to an AI, not a preserved record.

Kantian Principle: "Act only according to that maxim whereby you can at the same time will that it should become a universal law" – Kant. Posthumous data exploitation fails this test universally.

# Ethical Analysis: Virtue Ethics

07

## Honesty

LACKING

Platforms rarely disclose that responses are probabilistic hallucinations, not preserved memories. Users may believe they hear the authentic voice of the deceased.

## Responsibility

LACKING

Developers have not proactively assessed psychological harm risks. No mandatory grief counsellor collaboration or mental health safeguards are embedded in most platforms.

## Prudence

LACKING

Rushing to market without regulatory engagement, psychological research, or ethical frameworks demonstrates imprudence. StoryFile's institutional approach is a partial exception.

## Fairness

PARTIAL

HereAfter AI requires active family participation and consent, showing some fairness. However, commercial platforms like Project December place profit above user welfare.

## Compassion

MIXED

The intent to support grieving individuals demonstrates compassion. However, genuine compassion requires protecting long-term well-being, not short-term emotional comfort for profit.

## Integrity

LACKING

Microsoft's 2021 patent application for deceased-persona chatbots, filed without public ethical review, exemplifies a failure of corporate integrity and social responsibility.

# Ethical Analysis: Rights-Based Ethics

08

## Right to Privacy



Personal communications, voice, and likeness are harvested without clear posthumous consent frameworks. GDPR Art. 17 'right to be forgotten' is inapplicable to the dead in most EU member states (Harbinja, 2017).

## Right to Autonomy



The deceased cannot exercise agency over how their identity is represented. Family members speaking on their behalf hold no legally recognised proxy consent in most jurisdictions.

## Bereaved Person's Right to Healthy Grief



Prolonged simulated interaction may impede natural grieving. Mental health professionals note that griefbots can prevent acceptance – the final stage of the Kübler-Ross model.

## Right to Consent



No existing mechanism allows individuals to pre-authorise or refuse AI training on their posthumous data. Current ToS agreements are insufficient for this novel use case.

## Right to Identity



AI-generated responses may attribute false statements, beliefs, or expressions to the deceased. This constitutes a form of identity misrepresentation and post-mortem defamation.

## Data Rights of Next-of-Kin



Families who provide data have no guarantee of data deletion, ownership, or portability once the platform relationship ends. Profit from the deceased's data flows entirely to the corporation.

## Vulnerable Groups – Inadequately Protected

- Children who have lost a parent: AI simulation delays developmental adjustment to loss
- Elderly bereaved spouses: high risk of dependency and inability to distinguish AI from real
- Individuals with pre-existing mental health conditions: complicated grief risk significantly elevated
- Low-digital-literacy users: unable to critically assess probabilistic AI outputs vs real memory

## Human Well-Being & Relationships

- Noddings' ethics of care demands responsiveness to the actual emotional needs of the grieving, not simulated responsiveness
- Healthy grief requires acceptance of loss; griefbots can interrupt the relational transition needed to live without the deceased
- Care ethics values ongoing human relationships: counsellors, family networks, and community support outperform AI in authentic emotional responsiveness
- Platforms lack empathetic context awareness – the AI cannot detect when continued interaction is harmful to the user

**Noddings' Care Principle:** *Genuine care requires "engrossment" in the other's reality and "motivational displacement" toward their actual well-being. A commercial AI platform cannot substitute for genuine human care relationships in bereavement.*

## Missing Safeguards

- No mandatory mental health screening before account activation
- No usage-time limits or intervention triggers for at-risk users
- No requirement to offer professional grief counselling alongside the service
- No family consent protocols for the deceased's data

# Recommendations & Conclusion

10

1

**Ante-Mortem Consent Frameworks:** Mandate explicit pre-death consent mechanisms embedded in platform ToS, allowing individuals to authorise or prohibit posthumous AI training of their data.

2

**Mental Health Safeguards:** Require mandatory mental health risk screening, time-limited usage guidelines, and integration with licensed grief counsellors for all griefbot platforms.

3

**Posthumous Data Rights Legislation:** Governments must extend data protection laws (GDPR, PDPB) to cover the deceased, granting next-of-kin the right to access, correct, and erase AI-trained data.

4

**Algorithmic Transparency & Disclosure:** Platforms must clearly disclose AI-generated content at all interaction points. Users must know they are engaging with a probabilistic model, not a preserved consciousness.

5

**Independent Ethics Oversight:** Establish multi-disciplinary review boards (AI ethicists, grief counsellors, legal experts) to audit griefbot platforms before and after deployment.